

ECE 4750 Computer Architecture, Fall 2021

T01 Fundamental Processor Concepts

School of Electrical and Computer Engineering
Cornell University

revision: 2021-09-01-01-24

1	Instruction Set Architecture	3
1.1.	IBM 360 Instruction Set Architecture	5
1.2.	MIPS32 Instruction Set Architecture	7
1.3.	Tiny RISC-V Instruction Set Architecture	11
2	Processor Functional-Level Model	15
2.1.	Transactions and Steps	15
2.2.	TinyRV1 Simple Assembly Example	16
2.3.	TinyRV1 VVAdd Asm and C Program	17
2.4.	TinyRV1 Mystery Asm and C Program	18
3	Processor/Laundry Analogy	19
3.1.	Arch vs. μ Arch vs. VLSI Impl	19
3.2.	Processor Microarchitectural Design Patterns	20
3.3.	Transaction Diagrams	21
4	Analyzing Processor Performance	22

Copyright © 2016 Christopher Batten. All rights reserved. This handout was originally prepared by Prof. Christopher Batten at Cornell University for ECE 4750 / CS 4420. The handout has been modified and extended by Prof. Christina Delimitrou in 2017-2021. Download

and use of this handout is permitted for individual educational non-commercial purposes only. Redistribution either in part or in whole via both commercial or non-commercial means requires written permission.

1. Instruction Set Architecture

- By early 1960's, IBM had several incompatible lines of computers!
 - Defense : 701
 - Scientific : 704, 709, 7090, 7094
 - Business : 702, 705, 7080
 - Mid-Sized Business : 1400
 - Decimal Architectures : 7070, 7072, 7074
- Each system had its own:
 - Implementation and potentially even technology
 - Instruction set
 - I/O system and secondary storage (tapes, drums, disks)
 - Assemblers, compilers, libraries, etc
 - Application niche
- IBM 360 was the first line of machines to separate ISA from microarchitecture
 - Enabled same software to run on different current and future microarchitectures
 - Reduced impact of modifying the microarchitecture enabling rapid innovation in hardware

Application
Algorithm
Programming Language
Operating System
Instruction Set Architecture
Microarchitecture
Register-Transfer Level
Gate Level
Circuits
Devices
Physics

... the structure of a computer that a **machine language programmer** must understand to write a correct (**timing independent**) program for that machine.

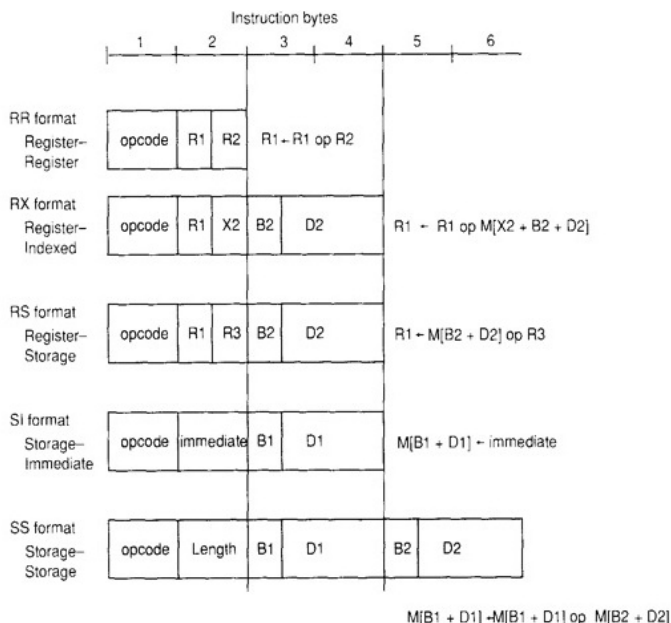
— Amdahl, Blaauw, Brooks, 1964

ISA is the contract between software and hardware

- 1. _____
 - Representations for characters, integers, floating-point
 - Integer formats can be signed or unsigned
 - Floating-point formats can be single- or double-precision
 - Byte addresses can ordered within a word as either little- or big-endian
- 2. _____
 - Registers: general-purpose, floating-point, control
 - Memory: different addresses spaces for heap, stack, I/O
- 3. _____
 - Register: operand stored in registers
 - Immediate: operand is an immediate in the instruction
 - Direct: address of operand in memory is stored in instruction
 - Register Indirect: address of operand in memory is stored in register
 - Displacement: register indirect, addr is added to immediate
 - Autoincrement/decrement: register indirect, addr is automatically adj
 - PC-Relative: displacement is added to the program counter
- 4. _____
 - Integer and floating-point arithmetic instructions
 - Register and memory data movement instructions
 - Control transfer instructions
 - System control instructions
- 5. _____
 - Opcode, addresses of operands and destination, next instruction
 - Variable length vs. fixed length

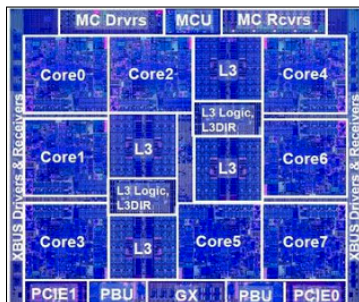
1.1. IBM 360 Instruction Set Architecture

- How is data represented?
 - 8-bit bytes, 16-bit half-words, 32-bit words, 64-bit double-words
 - IBM 360 is why bytes are 8-bits long today!
- Where can data be stored?
 - 2^{24} 32-bit memory locations
 - 16 general-purpose 32-bit registers and 4 floating-point 64-bit registers
 - Condition codes, control flags, program counter
- What operations can be done on data?
 - Large number of arithmetic, data movement, and control instructions



	Model 30	Model 70
Storage	8–64 KB	256–512 KB
Datapath	8-bit	64-bit
Circuit Delay	30 ns/level	5 ns/level
Local Store	Main store	Transistor registers
Control Store	Read only 1 μ s	Conventional circuits

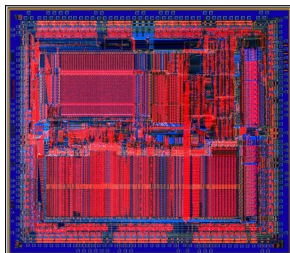
- IBM 360 instruction set architecture completely hid the underlying technological differences between various models
- **Significant Milestone: The first true ISA designed as a portable hardware-software interface**
- IBM 360: 50 years later ...
The zSeries z13 Microprocessor
 - 5 GHz in IBM 22 nm SOI
 - 4B transistors in 678 mm²
 - 17 metal layers
 - \approx 20K pads
 - Eight cores per chip
 - Aggressive out-of-order execution
 - Four-level cache hierarchy
 - On-chip 64MB eDRAM L3 cache
 - Off-chip 480MB eDRAM L4 cache
 - Can still run IBM 360 code!



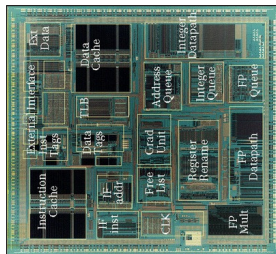
J. Warnock, et al., “22nm Next-Generation IBM System-Z Microprocessor,”
Int’l Solid-State Circuits Conference, Feb. 2016.

1.2. MIPS32 Instruction Set Architecture

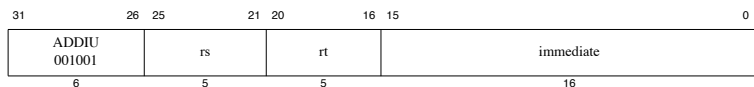
- How is data represented?
 - 8-bit bytes, 16-bit half-words, 32-bit words
 - 32-bit single-precision, 64-bit double-precision floating point
- Where can data be stored?
 - 2^{32} 32-bit memory locations
 - 32 general-purpose 32-bit registers, 32 SP (16 DP) floating-point registers
 - FP status register, Program counter
- How can data be accessed?
 - Register, immediate, displacement
- What operations can be done on data?
 - Large number of arithmetic, data movement, and control instructions
- How are instructions encoded?
 - Fixed-length 32-bit instructions



MIPS R2K: 1986, single-issue, in-order, off-chip caches, $2\text{ }\mu\text{m}$, 8–15 MHz, 110K transistors, 80 mm^2



MIPS R10K: 1996, quad-issue, out-of-order, on-chip caches, $0.35\text{ }\mu\text{m}$, 200 MHz, 6.8M transistors, 300 mm^2



Format: ADDIU *rt*, *rs*, *immediate*

MIPS32

Purpose: Add Immediate Unsigned Word

To add a constant to a 32-bit integer

Description: $GPR[rt] \leftarrow GPR[rs] + immediate$

The 16-bit signed *immediate* is added to the 32-bit value in GPR *rs* and the 32-bit arithmetic result is placed into GPR *rt*.

No Integer Overflow exception occurs under any circumstances.

Restrictions:

None

Operation:

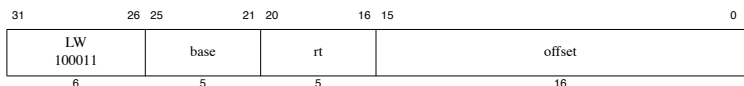
```
temp ← GPR[rs] + sign_extend(immediate)
GPR[rt] ← temp
```

Exceptions:

None

Programming Notes:

The term “unsigned” in the instruction name is a misnomer; this operation is 32-bit modulo arithmetic that does not trap on overflow. This instruction is appropriate for unsigned arithmetic, such as address arithmetic, or integer arithmetic environments that ignore overflow, such as C language arithmetic.



Format: LW *rt*, *offset*(*base*)

MIPS32

Purpose: Load Word

To load a word from memory as a signed value

Description: $GPR[rt] \leftarrow \text{memory}[GPR[base] + \text{offset}]$

The contents of the 32-bit word at the memory location specified by the aligned effective address are fetched, sign-extended to the GPR register length if necessary, and placed in GPR *rt*. The 16-bit signed *offset* is added to the contents of GPR *base* to form the effective address.

Restrictions:

The effective address must be naturally-aligned. If either of the 2 least-significant bits of the address is non-zero, an Address Error exception occurs.

Operation:

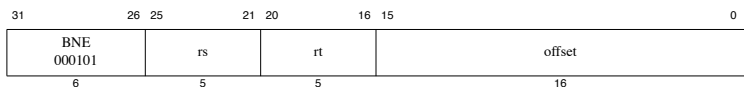
```

vAddr ← sign_extend(offset) + GPR[base]
if vAddr1..0 ≠ 02 then
    SignalException(AddressError)
endif
(pAddr, CCA) ← AddressTranslation (vAddr, DATA, LOAD)
memword ← LoadMemory (CCA, WORD, pAddr, vAddr, DATA)
GPR[rt] ← memword

```

Exceptions:

TLB Refill, TLB Invalid, Bus Error, Address Error, Watch



Format: BNE rs, rt, offset

MIPS32

Purpose: Branch on Not Equal

To compare GPRs then do a PC-relative conditional branch

Description: if $GPR[rs] \neq GPR[rt]$ then branch

An 18-bit signed offset (the 16-bit *offset* field shifted left 2 bits) is added to the address of the instruction following the branch (not the branch itself), in the branch delay slot, to form a PC-relative effective target address.

If the contents of GPR *rs* and GPR *rt* are not equal, branch to the effective target address after the instruction in the delay slot is executed.

Restrictions:

Processor operation is **UNPREDICTABLE** if a branch, jump, ERET, DERET, or WAIT instruction is placed in the delay slot of a branch or jump.

Operation:

```

I:    target_offset ← sign_extend(offset || 02)
        condition ← (GPR[rs] ≠ GPR[rt])
I+1: if condition then
        PC ← PC + target_offset
      endif
  
```

Exceptions:

None

Programming Notes:

With the 18-bit signed instruction offset, the conditional branch range is ± 128 KBytes. Use jump (J) or jump register (JR) instructions to branch to addresses outside this range.

1.3. Tiny RISC-V Instruction Set Architecture

- RISC-V instruction set architecture
 - Brand new free, open instruction set architecture
 - Significant excitement around RISC-V hardware/software ecosystem
 - Helping to energize “open-source hardware”
 - Specifically designed to encourage subsetting and extension
 - Link to official ISA manual on course webpage
- Tiny RISC-V instruction set architecture
 - Subset we use in this course
 - Small enough for teaching, powerful enough for running real C programs
 - How is data represented? _____
 - Where can data be stored? _____
 - How can data be accessed? _____
 - What ops can be done on data? _____
 - How are inst encoded? _____
 - <http://www.csl.cornell.edu/courses/ece4750/handouts>
- **TinyRV1:** Small subset suitable for lecture, homeworks, exams
 - _____
 - _____
 - _____
- **TinyRV2:** Subset suitable for lab assignments and capable of executing simple C programs without an operating system
 - add, addi, sub, mul, and, andi, or, ori, xor, xori
 - slt, slti, sltu, sltiu
 - sra, srai, srl, srli, sll, slli
 - lui, auipc, lw, sw
 - jal, jalr, beq, bne, blt, bge, bltu, bgeu
 - csrr, csrw

TinyRV1 instruction assembly, semantics, and encoding

ADD

add rd, rs1, rs2

$R[rd] \leftarrow R[rs1] + R[rs2]$

$PC \leftarrow PC + 4$

31	25	24	20	19	15	14	12	11	7	6	0
0000000	rs2	rs1	000	rd	0110011						

ADDI

addi rd, rs1, imm

$R[rd] \leftarrow R[rs1] + \text{sext}(imm)$

$PC \leftarrow PC + 4$

31	20	19	15	14	12	11	7	6	0
imm		rs1		000		rd		0010011	

MUL

mul rd, rs1, rs2

$R[rd] \leftarrow R[rs1] \times R[rs2]$

$PC \leftarrow PC + 4$

31	25	24	20	19	15	14	12	11	7	6	0
0000001	rs2	rs1	000	rd	0110011						

LW

lw rd, imm(rs1)

$R[rd] \leftarrow M[R[rs1] + \text{sext}(imm)]$

$PC \leftarrow PC + 4$

31	20	19	15	14	12	11	7	6	0
imm		rs1	010		rd		0000011		

SW

sw rs2, imm(rs1)

$M[R[rs1] + \text{sext}(imm)] \leftarrow R[rs2]$

$PC \leftarrow PC + 4$

31	25	24	20	19	15	14	12	11	7	6	0
imm	rs2	rs1	010	imm	0100011						

$imm = \{ inst[31:25], inst[11:7] \}$

JAL

jal rd, imm

$R[rd] \leftarrow PC + 4$

$PC \leftarrow PC + \text{sext}(imm)$

31	12	11	7	6	0
imm			rd	1101111	

$imm = \{ inst[31], inst[19:12], inst[20], inst[30:21], 0 \}$

JR

jr rs1

$PC \leftarrow R[rs1]$

31	20	19	15	14	12	11	7	6	0
000000000000		rs1		000		00000		1100111	

BNE

bne rs1, rs2, imm

if ($R[rs1] \neq R[rs2]$) $PC \leftarrow PC + \text{sext}(imm)$

else $PC \leftarrow PC + 4$

31	25	24	20	19	15	14	12	11	7	6	0
imm	rs2	rs1	001	imm	1100011						

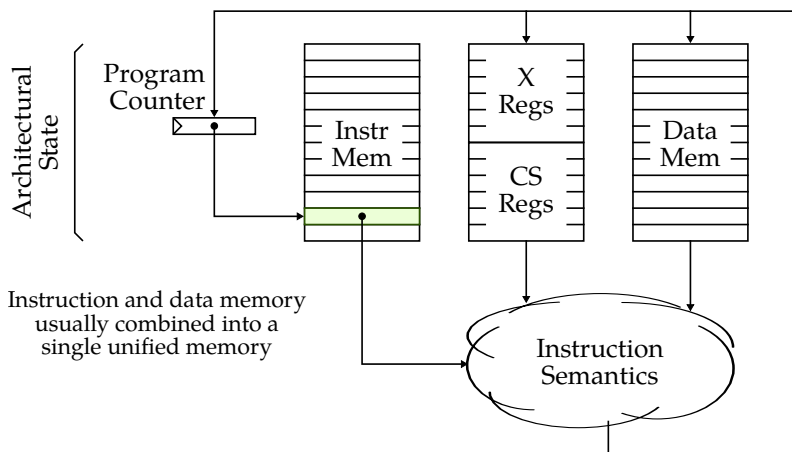
$imm = \{ inst[31], inst[7], inst[30:25], inst[11:8], 0 \}$

Base Integer Instructions: RV32I, RV64I, and RV128I						RV Privileged Instructions						
Category	Name	Fmt	RV32I Base		+RV(64,128)	Category	Name	RV mnemonic				
Loads	Load Byte	I	LB	rd,rs1,imm	L{D Q}	rd,rs1,imm	CSR Access	Atomic R/W	CSRWR rd,csr,rs1			
	Load Halfword	I	LH	rd,rs1,imm				Atomic Read & Set Bit	CSRRS rd,csr,rs1			
	Load Word	I	LW	rd,rs1,imm				Atomic Read & Clear Bit	CSRRC rd,csr,rs1			
	Load Byte Unsigned	I	LBU	rd,rs1,imm				Atomic R/W Imm	CSRRIW rd,csr,imm			
	Load Half Unsigned	I	LHU	rd,rs1,imm				L{W D}U rd,rs1,imm	Atomic Read & Set Bit Imm	CSRRII rd,csr,imm		
Stores	Store Byte	S	SB	rs1,rs2,imm	S{D Q}	rs1,rs2,imm	Atomic Read & Clear Bit Imm	CSRRCI rd,csr,imm				
	Store Halfword	S	SH	rs1,rs2,imm								
	Store Word	S	SW	rs1,rs2,imm								
Shifts	Shift Left	R	SLL	rd,rs1,rs2	SLL{W D}	rd,rs1,rs2	Change Level	Env. Call	ECALL			
	Shift Left Immediate	I	SLLI	rd,rs1,shamt	SLLI{W D}	rd,rs1,shamt		Environment Breakpoint	EBREAK			
	Shift Right	R	SRL	rd,rs1,rs2	SRL{W D}	rd,rs1,rs2		Environment Return	ERET			
	Shift Right Immediate	I	SRLI	rd,rs1,shamt	SRLI{W D}	rd,rs1,shamt		Trap Redirect to Supervisor	MRTS			
	Shift Right Arithmetic	R	SRA	rd,rs1,rs2	SRA{W D}	rd,rs1,rs2			Redirect Trap to Hypervisor	MRTS		
Arithmetic	Shift Right Arith Imm	I	SRAI	rd,rs1,shamt	SRAI{W D}	rd,rs1,shamt	Hypervisor Trap to Supervisor	HRTS				
	ADD	R	ADD	rd,rs1,rs2	ADD{W D}	rd,rs1,rs2	Interrupt	Wait for Interrupt	WFI			
	ADD Immediate	I	ADDI	rd,rs1,imm	ADDI{W D}	rd,rs1,imm		MMU	Supervisor FENCE	SFENCE.VM rs1		
	SUBtract	R	SUB	rd,rs1,rs2	SUB{W D}	rd,rs1,rs2						
	Load Upper Imm	U	LUI	rd,imm								
Add Upper Imm to PC	U	AUIPC	rd,imm									
Optional Compressed (16-bit) Instruction Extension: RVC												
Category	Name	Fmt	RVC		RVI equivalent							
Logical	XOR	R	XOR	rd,rs1,rs2	CL	C.LW	rd',rs1',imm	LW	rd',rs1',imm*4			
	XOR Immediate	I	XORI	rd,rs1,imm		C.LWSP	rd,imm	LW	rd,sp,imm*4			
	OR	R	OR	rd,rs1,rs2		C.LD	rd',rs1',imm	LD	rd',rs1',imm*8			
	OR Immediate	I	ORI	rd,rs1,imm		C.LDSP	rd,imm	LD	rd,sp,imm*8			
	AND	R	AND	rd,rs1,rs2		C.LQ	rd',rs1',imm	LQ	rd',rs1',imm*16			
Compare	AND Immediate	I	ANDI	rd,rs1,imm	CL	C.LQSP	rd,imm	LQ	rd,sp,imm*16			
	Set <	R	SLT	rd,rs1,rs2		CS	C.SW	rs1',rs2',imm	SW	rs1',rs2',imm*4		
	Set < Immediate	I	SLTI	rd,rs1,imm			C.SWSP	rs2,imm	SW	rs2,sp,imm*4		
	Set < Unsigned	R	SLTU	rd,rs1,rs2			C.SD	rs1',rs2',imm	SD	rs1',rs2',imm*8		
	Set < Imm Unsigned	I	SLTIU	rd,rs1,imm			C.SDSP	rs2,imm	SD	rs2,sp,imm*8		
Branches	Branch =	SB	BEQ	rs1,rs2,imm	CS		C.SQ	rs1',rs2',imm	SQ	rs1',rs2',imm*16		
	Branch #	SB	BNE	rs1,rs2,imm		CSS	C.SQSP	rs2,imm	SQ	rs2,sp,imm*16		
	Branch <	SB	BLT	rs1,rs2,imm			CR	C.ADD	rd,rs1	ADD	rd,rd,rs1	
	Branch >	SB	BGE	rs1,rs2,imm				C.ADDW	rd,rs1	ADDW	rd,rd,imm	
	Branch < Unsigned	SB	BLTU	rs1,rs2,imm				C.ADDI	rd,rs1	ADDI	rd,rd,imm	
Jump & Link	Branch > Unsigned	SB	BGEU	rs1,rs2,imm	CI			C.ADDIW	rd,imm	ADDIW	rd,rd,imm	
	J&L	UJ	JAL	rd,imm		CIW		C.ADDI16SP	x0,imm	ADDI	sp,sp,imm*16	
	Jump & Link Register	UJ	JALR	rd,rs1,imm			C	C.ADDI4SPN	rd',imm	ADDI	rd',sp,imm*4	
	Synch	I	FENCE					C.LI	rd,imm	ADDI	rd,x0,imm	
	Synch Instr & Data	I	FENCE.I					C.LUI	rd,imm	LUI	rd,imm	
System	System CALL	I	SCALL		CR			C.MV	rd,rs1	LD	rd,rs1,x0	
	System BREAK	I	SBREAK			CR		C.SUB	rd,rs1	SUB	rd,rd,rs1	
	Counters Read CYCLE	I	RDYCYCLE	rd			C	C.SLLI	rd,imm	SLLI	rd,rd,imm	
	Read CYCLE upper Half	I	RDYCYCLEH	rd				CB	C.BEQZ	rs1',imm	BEQ	rs1',x0,imm
	Read TIME	I	RDTIME	rd					CB	C.BNEZ	rs1',imm	BNE
Read TIME upper Half	I	RDTIMEH	rd	C	C.J					imm	JAL	x0,imm
Read INSTR RETired	I	RDINSTRET	rd		CR	C.JR				rd,rs1	JALR	x0,rs1,0
Read INSTR upper Half	I	RDINSTRETH	rd			CJ	C.JAL			imm	JAL	ra,imm
Counters	Read INSTR upper Half	I	RDINSTRETH				rd	CR		C.JALR	rs1	JALR
	System	Env. BREAK	CI				FRRPAK			CI	FRRPAK	

Optional Multiply-Divide Instruction Extension: RVM							
Category	Name	Fmt	RV32M (Multiply-Divide)	+RV{64,128}			
Multiply	Multiply	R	MUL rd,rs1,rs2	MUL{W D} rd,rs1,rs2			
	Multiply upper Half	R	MULH rd,rs1,rs2				
	Multiply Half Sign/Uns	R	MULHSU rd,rs1,rs2				
	Multiply upper Half Uns	R	MULHU rd,rs1,rs2				
Divide	DIVide	R	DIV rd,rs1,rs2	DIV{W D} rd,rs1,rs2			
	DIVide Unsigned	R	DIVU rd,rs1,rs2				
Remainder	REMAinder	R	REM rd,rs1,rs2	REM{W D} rd,rs1,rs2			
	REMAinder Unsigned	R	REMU rd,rs1,rs2	REMU{W D} rd,rs1,rs2			
Optional Atomic Instruction Extension: RVA							
Category	Name	Fmt	RV32A (Atomic)	+RV{64,128}			
Load	Load Reserved	R	LR.W rd,rs1	LR.{D Q} rd,rs1			
Store	Store Conditional	R	SC.W rd,rs1,rs2	SC.{D Q} rd,rs1,rs2			
Swap	SWAP	R	AMOSWAP.W rd,rs1,rs2	AMOSWAP.{D Q} rd,rs1,rs2			
Add	ADD	R	AMOADD.W rd,rs1,rs2	AMOADD.{D Q} rd,rs1,rs2			
Logical	XOR	R	AMOXOR.W rd,rs1,rs2	AMOXOR.{D Q} rd,rs1,rs2			
	AND	R	AMOAND.W rd,rs1,rs2	AMOAND.{D Q} rd,rs1,rs2			
	OR	R	AMOOR.W rd,rs1,rs2	AMOOR.{D Q} rd,rs1,rs2			
Min/Max	MINimum	R	AMOMIN.W rd,rs1,rs2	AMOMIN.{D Q} rd,rs1,rs2			
	MAXimum	R	AMOMAX.W rd,rs1,rs2	AMOMAX.{D Q} rd,rs1,rs2			
	MINimum Unsigned	R	AMOMINU.W rd,rs1,rs2	AMOMINU.{D Q} rd,rs1,rs2			
	MAXimum Unsigned	R	AMOMAXU.W rd,rs1,rs2	AMOMAXU.{D Q} rd,rs1,rs2			
Three Optional Floating-Point Instruction Extensions: RVF, RVD, & RVQ							
Category	Name	Fmt	RV32{F D Q} (HP/SP,DP,QP,FP Pt)	+RV{64,128}			
Move	Move from Integer	R	FMV.{H S}.X rd,rs1	FMV.{D Q}.X rd,rs1			
	Move to Integer	R	FMV.X.{H S} rd,rs1	FMV.X.{D Q} rd,rs1			
Convert	Convert from Int	R	FCVT.{H S D Q}.W rd,rs1	FCVT.{H S D Q}.L{L T} rd,rs1			
	Convert from Int Unsigned	R	FCVT.{H S D Q}.WU rd,rs1	FCVT.{H S D Q}.L{L T}U rd,rs1			
	Convert to Int	R	FCVT.W.{H S D Q} rd,rs1	FCVT.L{L T}.H{S D Q} rd,rs1			
	Convert to Int Unsigned	R	FCVT.WU.{H S D Q} rd,rs1	FCVT.L{L T}U.H{S D Q} rd,rs1			
Load	Load	I	FLW.{D,Q} rd,rs1,imm				
Store	Store	S	FSW.{D,Q} rd,rs1,rs2,imm				
			RISC-V Calling Convention				
Arithmetic			Register	ABI Name	Saver	Description	
	ADD	R	FADD.{S D Q} rd,rs1,rs2	x0	zero	---	Hard-wired zero
	SUBtract	R	FSUB.{S D Q} rd,rs1,rs2	x1	ra	Caller	Return address
	MULTiply	R	FMUL.{S D Q} rd,rs1,rs2	x2	sp	Callee	Stack pointer
	DIVide	R	FDIV.{S D Q} rd,rs1,rs2	x3	gp	---	Global pointer
	SQuare RooT	R	FSQRT.{S D Q} rd,rs1	x4	tp	---	Thread pointer
Mul-Add	Multiply-ADD	R	FMAADD.{S D Q} rd,rs1,rs2,rs3	x5-7	t0-2	Caller	Temporaries
	Multiply-SUBtract	R	FMSUB.{S D Q} rd,rs1,rs2,rs3	x8	s0/fp	Callee	Saved register/frame pointer
	Negative Multiply-SUBtract	R	FNMSUB.{S D Q} rd,rs1,rs2,rs3	x9	s1	Callee	Saved register
	Negative Multiply-ADD	R	FNMAADD.{S D Q} rd,rs1,rs2,rs3	x10-11	a0-1	Caller	Function arguments/return values
Sign Inject	SIGN source	R	FSGNJ.{S D Q} rd,rs1,rs2	x12-17	a2-7	Caller	Function arguments
	Negative SIGN source	R	FSGNJN.{S D Q} rd,rs1,rs2	x18-27	s2-11	Callee	Saved registers
	Xor SIGN source	R	FSGNJX.{S D Q} rd,rs1,rs2	x28-31	t3-t6	Caller	Temporaries
Min/Max	MINimum	R	FMIN.{S D Q} rd,rs1,rs2	f0-7	ft0-7	Caller	FP temporaries
	MAXimum	R	FMAX.{S D Q} rd,rs1,rs2	f8-9	fs0-1	Callee	FP saved registers
Compare	Compare Float =	R	FEQ.{S D Q} rd,rs1,rs2	f10-11	fa0-1	Caller	FP arguments/return values
	Compare Float <	R	FLT.{S D Q} rd,rs1,rs2	f12-17	fa2-7	Caller	FP arguments
	Compare Float ≤	R	FLE.{S D Q} rd,rs1,rs2	f18-27	fs2-11	Callee	FP saved registers
Categorization	Classify Type	R	FCLASS.{S D Q} rd,rs1	f28-31	ft8-11	Caller	FP temporaries
Configuration	Read Status	R	FRCSR rd				
	Read Rounding Mode	R	FRRM rd				
	Read Flags	R	FRFLAGS rd				
	Swap Status Reg	R	FRCSR rd,rs1				
	Swap Rounding Mode	R	FRRM rd,rs1				
	Swap Flags	R	FRFLAGS rd,rs1				
	Swap Rounding Mode Imm	I	FRSRI rd,imm				
	Swap Flags Imm	I	FRFLSGI rd,imm				

RISC-V calling convention and five optional extensions: 10 multiply-divide instructions (RV32M); 11 optional atomic instructions (RV32A); and 25 floating-point instructions each for single-, double-, and quadruple-precision (RV32F, RV32D, RV32Q). The latter add registers f0-f31, whose width matches the widest precision, and a floating-point control and status register fcsr. Each larger address adds some instructions: 4 for RVM, 11 for RVA, and 6 each for RVF/D/Q. Using regex notation, {} means set, so L{D|Q} is both LD and LQ. See risc.org. (8/21/15 revision)

2. Processor Functional-Level Model



2.1. Transactions and Steps

- We can think of each instruction as a **transaction**
- Executing a transaction involves a sequence of **steps**

	add	addi	mul	lw	sw	jal	jr	bne
Fetch Instruction								
Decode Instruction								
Read Registers								
Register Arithmetic								
Read Memory								
Write Memory								
Write Registers								
Update PC								

2.2. TinyRV1 Simple Assembly Example

Static Asm Sequence	Instruction Semantics
loop: lw x1, 0(x2)	
add x3, x3, x1	
addi x2, x2, 4	
bne x1, x0, loop	

Worksheet illustrating processor functional-level model

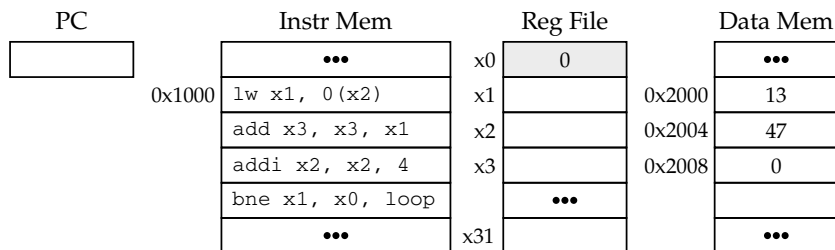


Table illustrating processor functional-level model

PC	Dynamic Asm Sequence	x1	x2	x3
	lw x1, 0(x2)			
	add x3, x3, x1			
	addi x2, x2, 4			
	bne x1, x0, loop			
	lw x1, 0(x2)			
	add x3, x3, x1			

2.3. TinyRV1 Vector-Vector Add Assembly and C Program

C code for doing element-wise vector addition.

Equivalent TinyRV1 assembly code. Arguments are passed in x12–x17, return value is stored to x10, return address is stored in x1, and temporaries are stored in x5–x7.

Note that we are ignoring the fact that our assembly code will not function correctly if $n \leq 0$. Our assembly code would need an additional check before entering the loop to ensure that $n > 0$. Unless otherwise stated, we will assume in this course that array bounds are greater than zero to simplify our analysis.

2.4. TinyRV1 Mystery Assembly and C Program

What is the C code corresponding to the TinyRV1 assembly shown below? Assume assembly implements a function.

```
addi x5, x0, 0
```

```
loop:
```

```
lw    x6, 0(x12)
```

```
bne   x6, x14, foo
```

```
addi  x10, x5, 0
```

```
jr    x1
```

```
foo:
```

```
addi  x12, x12, 4
```

```
addi  x5, x5, 1
```

```
bne   x5, x13, loop
```

```
addi  x10, x0, -1
```

```
jr    x1
```

3. Processor/Laundry Analogy

• Processor

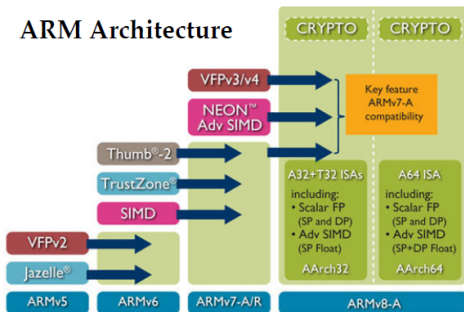
- Instructions are “transactions” that execute on a processor
- Architecture: defines the hardware/software interface
- Microarchitecture: how hardware executes sequence of instructions

• Laundry

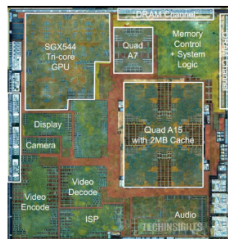
- Cleaning a load of laundry is a “transaction”
- Architecture: high-level specification, dirty clothes in, clean clothes out
- Microarchitecture: how laundry room actually processes multiple loads

3.1. Arch vs. μ Arch vs. VLSI Impl

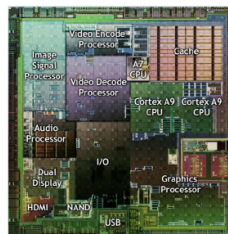
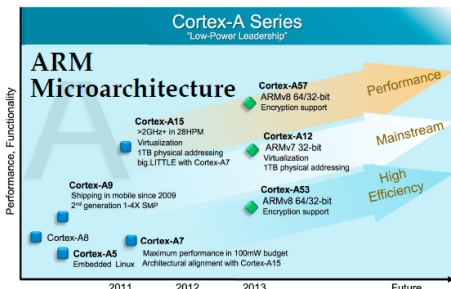
ARM Architecture



ARM VLSI Implementation



Samsung Exynos Octa



NVIDIA Tegra 2

3.2. Processor Microarchitectural Design Patterns

Transaction Steps



Washing
(30 min)



Drying
(30 min)



Folding
(30 min)



Storing
(30 min)

Four Types of Transactions

0 hr 1 hr 2 hr Transaction Latency

Anne's Load



2.0 hr

Anne requires all four steps

Ben's Load



1.0 hr

Ben is messy, leaves unfolded clothes in his laundry basket

Cathy's Load



1.5 hr

Cathy does not have a bureau, leaves folded clothes in basket

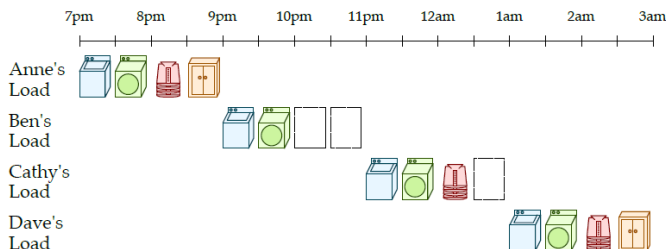
Dave's Load



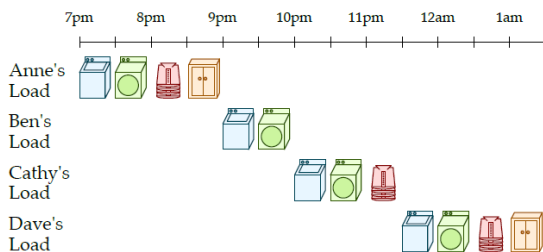
2.0 hr

Dave requires all four steps

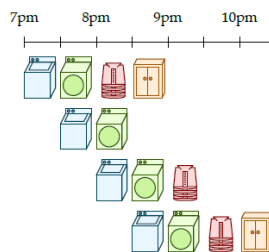
Fixed Time Slot Laundry (Single-Cycle Processors)



Variable Time Slot Laundry (FSM Processors)



Pipelined Laundry



3.3. Transaction Diagrams



W: Washing



D: Drying



F: Folding



S: Storing

[illegible][illegible]

Key Concepts

- **Transaction latency** is the time to complete a single transaction
- **Execution time** or **total latency** is the time to complete a sequence of transactions
- **Throughput** is the number of transactions executed per unit time

4. Analyzing Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Avg Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

- Instructions / program depends on source code, compiler, ISA
- Avg cycles / instruction (CPI) depends on ISA, microarchitecture
- Time / cycle depends upon microarchitecture and implementation

Using our first-order equation for processor performance and a functional-level model, the execution time is just the number of dynamic instructions.

Microarchitecture	CPI	Cycle Time
Single-Cycle Processor	1	long
FSM Processor	>1	short
Pipelined Processor	≈1	short



Students often confuse “Cycle Time” with the execution time of a sequence of transactions measured in cycles.
“Cycle Time” is the clock period or the inverse of the clock frequency.

Estimating dynamic instruction count

Estimate the dynamic instruction count for the vector-vector add example assuming n is 64?

```
loop:
    lw    x5,  0(x13)
    lw    x6,  0(x14)
    add   x7,  x5,  x6
    sw    x7,  0(x12)
    addi  x13, x12, 4
    addi  x14, x14, 4
    addi  x12, x12, 4
    addi  x15, x15, -1
    bne   x15, x0,  loop
    jr    x1
```

Estimate the dynamic instruction count for the mystery program assuming n is 64 and that we find a match on the final element.

```
    addi  x5,  x0,  0
loop:
    lw    x6,  0(x12)
    bne   x6,  x14, foo
    addi  x10, x5,  0
    jr    x1
foo:
    addi  x12, x12, 4
    addi  x5,  x5,  1
    bne   x5,  x13, loop
    addi  x10, x0,  -1
    jr    x1
```